BridgeFinder
Finding Research Collaborators Across Communities at UIUC

Evan Lin Eric Liu
University of Illinois Urbana-Champaign

CS 470 Mid-Semester Report, October 2026

Abstract

Students who want research experience often know an instructor but few researchers outside that instructor's circle. BridgeFinder helps them find related researchers and see who might connect them. We have built a three-screen prototype, analyzed UIUC publications from 2015 to 2024, and gathered feedback from five students. Our historical test finds that shared-contact ranking predicts more later collaborations than our method. The next study will test whether students nevertheless find our suggestions relevant, unfamiliar, and useful for contacting someone.

1 Objective

A student who enjoyed a computing class might want to apply those skills in biology or agriculture. Searching for papers can reveal relevant names, but it does not tell the student whom they could approach or how to start a conversation.

BridgeFinder helps UIUC undergraduate and early-stage graduate students explore research beyond their existing contacts. A student chooses a familiar researcher, sees suggestions outside that person's collaboration group, and opens a card with topics, papers, and a possible introduction route. The starting researcher's work stands in for the student's interests, so the student needs no publications of their own.

The challenge is to suggest someone relevant who also extends the student's usual circle. The prototype supports this sequence. We still need to test whether the suggestions help students choose someone to contact.

2 Related Work

Researcher search. Illinois Experts [12] collects institutional profiles and publications, making expertise easy to check. Google Scholar [9] helps readers find papers and follow authors. These tools help students identify researchers. BridgeFinder also shows how a suggested researcher connects to someone the student knows.

Literature maps. Connected Papers [10] makes related papers visible in a local map. ResearchRabbit [11] supports continued discovery from paper collections. Both make exploration easier, but their main objects are papers. BridgeFinder adapts visual exploration to people and possible introductions.

Recommendations. Liben-Nowell and Kleinberg [5] compare network-based link predictions. Shared-contact measures are understandable, but favor nearby people. We compare with simple triadic closure, which ranks candidates by how many co-authors they share with the starting researcher. Tang et al. [8] combine research interests and relationships for cross-domain collaboration recommendation. That supports considering both topic and network evidence, but ranking quality alone does not show whether a student can use a suggestion. We therefore also plan a student evaluation.

Network theory. Granovetter [4] explains how weak ties can reach beyond close circles. Burt [2] explains the advantage of bridging separate groups. We use these ideas to look beyond a familiar collaboration group. A co-authorship path alone cannot tell us whether the people know each other well enough to make an introduction.

3 Course Concepts and Network Analysis

3.1 Graph and communities

Brokerage involves connecting people across groups that have few ties to one another. We use this course concept to suggest researchers outside a familiar person's collaboration group. We use OpenAlex [7] publications from 2015–2024. Nodes are authors affiliated with UIUC on a collected paper. An edge joins two authors who wrote together, weighted by their number of joint papers. We retain papers with at least two UIUC authors and omit external collaborators.

Our dataset contains 49,651 papers that meet this rule, retained from an exhausted harvest of 102,947 OpenAlex works matching our query. OpenAlex is mutable and does not index every real-world paper. Table 1 describes the 2015–2019 graph used to make recommendations.

Table 1. Co-authorship graphs from 2015–2019 papers.

Measure	Dataset
Authors	19,389
Edges	78,281
Communities	842
Modularity	0.884

We use Louvain [1] to find communities with many collaborations inside each group. It seeks high *modularity*, meaning more within-group edge weight than expected from node strengths. We chose it over Girvan–Newman [3], whose repeated shortest-path calculations are costly on these graphs.

The high modularity suggests concentrated collaboration groups, making cross-group discovery worth exploring. Disconnected components also raise modularity, so this is not evidence of disciplinary separation. Communities are numbered groups, not verified departments.

3.2 Choosing suggestions

Following the proposal, candidates must be in another community and at least three hops from the starting researcher. One hop is one co-authorship edge. This excludes existing collaborators and people who already share a co-author. Candidates must also be reachable so we can show a route. All 500 displayed suggestions across the demo's 50 starting researchers have three-edge routes.

The ranking combines topic overlap with a bonus for introduction paths, as shown below.

$$\text{Score}(u, v) = S(u, v) + 0.5\frac{c_3(u, v)}{1 + c_3(u, v)}. \qquad (1)$$

S is the number of shared topics divided by the number of distinct topics across both researchers. Each researcher has up to ten topics from their 2015–2019 papers. c_3 counts three-edge paths through two distinct intermediaries. The bonus grows with the number of routes but stays below 0.5. More distant candidates remain eligible without that bonus. The coefficient and saturation are fixed implementation choices, not learned probabilities.

Without the network, the app would lose its group boundary, distance rule, and introduction routes and become a topic search tool. The resulting suggestions extend beyond shared contacts, although the graph does not tell us whether an introduction would be welcome.

4 Design and Implementation

The prototype takes the student through three views.

1. Researchers. The student searches 50 starting researchers by name or topic. We chose highly connected authors across large communities so the demo has enough suggestions. The landing page explains the steps and defines a researcher card. As in profile search [12], it provides an entry point without requiring a student profile.

2. Network. A map shows the selected researcher, suggestions, and intermediaries. Solid lines show observed co-authorships. Orange dashed lines show suggested connections. This borrows visual exploration from literature maps [10, 11] while focusing on people.

3. Suggestions. Selecting a name or dashed line opens a card with an observed shortest route, topics, and example papers. The route appears first in response to student feedback. An optional explanation shows the ranking score. Students can edit, copy, or download an outreach draft. The app does not send email. Figure 1 shows the map and route detail.

The app, historical test, and study cards share one ranking implementation. Recommendations are computed in Python and displayed in D3. The core flow works on desktop and mobile, including keyboard navigation, and has been browser-tested for all 50 starting researchers.

5 Users and Feedback

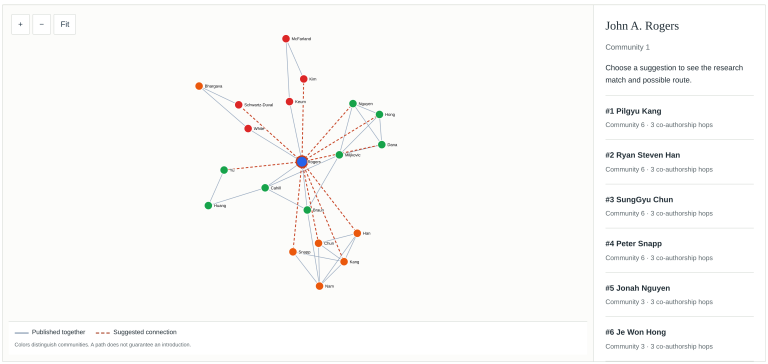
Our target users want research opportunities but have few contacts beyond their classes or department. We gathered feedback in five individual, 30-minute sessions in which students talked through their use of the prototype. Participants were three CS/CE undergraduate juniors or seniors and two graduate students from Bioengineering and Information Sciences.

The undergraduates fit the entry problem of exploring research before having publications. The graduate students add perspectives from application areas. Five sessions suit an initial search for interaction problems [6], rather than a campus-wide usability estimate. We made four changes in response to the feedback.

Getting started. One student trying to find a relevant professor did not understand that they first needed to select a researcher to enter the graph. We added instructions on the landing page explaining that students should choose a starting researcher, explore a suggestion, and then review its route. The page now explains that the starting person's publications guide the suggestions.

Understanding the interface. Students found the interface unintuitive, prompting an overhaul with simpler layouts and clearly named views. They were also unsure what a card meant. The landing page now defines a researcher card as a profile with topics, example papers, and a possible introduction route. The map legend distinguishes observed collaborations from suggestions.

Readable map. The graph showed too many nodes and was difficult to follow. We narrowed the display to the selected researcher, suggestions, and people along their routes. Community colors help distinguish groups without showing the entire campus graph.



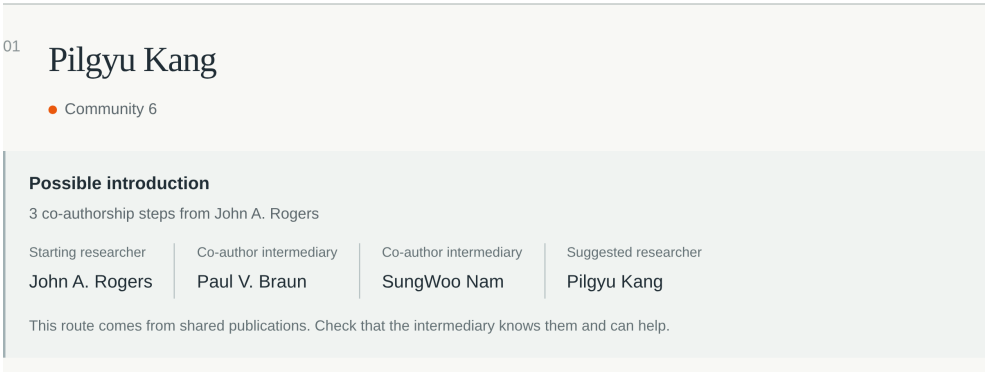


Figure 1. The upper image shows John A. Rogers's map. Solid lines represent observed co-authorships, and orange dashed lines represent suggestions. The lower image shows a detail of Pilgyu Kang's card, with a route through Paul V. Braun and SungWoo Nam. Topics and papers appear below this detail. The route is observed in publications, not a confirmed introduction.

Correct labels. Some tags in the original design did not match the researchers shown. We corrected the profile labeling and distinguish inferred research areas from collaboration groups. These labels do not claim verified departmental appointments.

The new entry instructions and card definition have not yet been retested with students. The observations explain why we made the revisions. Further testing is needed to find out whether they help. We do not report issue counts or participant quotations.

6 Evaluation

6.1 Historical prediction

We used 2015–2019 papers to build the graph and topic profiles, then checked whether suggestions became new cross-community co-authors in 2020–2024. We evaluated 3,492 starting researchers in this dataset. Each had at least one new cross-community tie between training-graph authors. Researchers without such future ties are excluded.

We compare with **Topic-Only** on the same candidate pool and **Triadic closure**, the shared-contact baseline in our proposal. We rank these candidates by their number of shared co-authors, with no additional weighting. These candidates are two hops away. Restricting that baseline to our three-or-more-hop pool would give only zero scores. All methods predict the same future ties, so this compares discovery strategies with different pools.

We measure the fraction of each researcher's top ten suggestions that became collaborators in the later period. This is precision at ten, or **P@10**. It matches the proposal's test of whether the shortlist includes people who later worked together. We average the fraction across starting researchers.

Table 2. Prediction of new 2020–2024 cross-community ties. Higher is better.

Method	P@10
Topic-Only	0.0079
Triadic closure	0.0287
BridgeFinder	0.0131

Table 2 shows that BridgeFinder predicts more later ties than Topic-Only (0.0131 vs 0.0079), but fewer than Triadic closure (0.0287). It therefore does not meet the proposal's goal of outperforming both baselines. We have not estimated uncertainty, and these comparisons do not isolate distance from scoring. The student study will test whether the suggestions are useful for discovery despite lower historical prediction than triadic closure.

6.2 Planned student study

We will recruit 15 students, aiming for three each from CS, ECE, Bioengineering, Crop Sciences, and Information Sciences. This adds backgrounds beyond those in our first feedback round, but the study remains exploratory.

Each student will choose a starting researcher, open a suggestion and outreach draft, and explain the route and dashed line without coaching. A facilitator will record task completion and whether the student identifies the intermediaries and distinguishes suggestions from existing ties.

For the card comparison, students will choose one of five starting researchers and rate up to 15 shuffled cards, with five from each method above. Method labels are hidden. Shared candidates appear once and count for every method that selected them. We will collect the following measures.

- **Relevant discovery.** Students will rate relevance to the starting researcher's work and novelty to themselves on scales from 1 to 7. For BridgeFinder, we target means of at least 5 and a 25-percentage-point gain over triadic closure in cards rated at least 5 on both. Topic-Only is a second comparison.
- **Willingness to contact.** Students will rate how realistic contact feels given the route on a scale from 1 to 7. For BridgeFinder, we target a mean of at least 5 and improvement over Topic-Only. Because the methods suggest different people, this rating cannot isolate the effect of showing a path.
- **Usability.** At least 12 of 15 students should complete the task and explain the line types correctly without help. The ten-item System Usability Scale will assess the interactive app, with a target mean of 80/100.

We will compare methods using participant averages, not treat individual cards as independent responses. The rating and analysis tools are implemented. We have not yet run the study.

7 Limitations and Next Steps

Co-authorship cannot confirm an introduction or mentoring availability. Our dataset omits external and informal ties, and OpenAlex does not capture every publication. The demo offers 50 starting researchers, and the card study offers five, so some students may not find a suitable starting point. A researcher's topics may also be a poor proxy for a student's interests.

In November, we will run the student study and revise the interface. Any ranking changes will use separate validation data. We will use the findings to assess whether the suggestions help students identify researchers they would consider contacting.

Contributions and AI Use

Evan Lin worked on data collection, prediction tests, and the map. Eric Liu worked on scoring, citation analysis, study tools, and formative sessions. Both contributed to the report. Claude assisted with early development and drafting. OpenAI Codex assisted with checks, tests, and revision. No participant responses or quotations were generated.

References

[1] V. D. Blondel, J.-L. Guillaume, R. Lambiotte, and E. Lefebvre. Fast unfolding of communities in large networks. *Journal of Statistical Mechanics: Theory and Experiment*, 2008(10):P10008, 2008.

[2] R. S. Burt. *Structural Holes: The Social Structure of Competition*. Harvard University Press, Cambridge, MA, 1992.

[3] M. Girvan and M. E. J. Newman. Community structure in social and biological networks. *Proceedings of the National Academy of Sciences*, 99(12):7821–7826, 2002.

[4] M. S. Granovetter. The strength of weak ties. *American Journal of Sociology*, 78(6):1360–1380, 1973.

[5] D. Liben-Nowell and J. Kleinberg. The link-prediction problem for social networks. *Journal of the American Society for Information Science and Technology*, 58(7):1019–1031, 2007.

[6] J. Nielsen. *Usability Engineering*. Morgan Kaufmann, San Francisco, CA, 1993.

[7] J. Priem, H. Piwowar, and R. Orr. Openalex: A fully-open index of scholarly works, authors, venues, institutions, and concepts. *arXiv preprint arXiv:2205.01833*, 2022.

[8] J. Tang, S. Wu, J. Sun, and H. Su. Cross-domain collaboration recommendation. In *Proceedings of the 18th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '12)*, pages 1285–1293, 2012.

[9] Google. Google Scholar: About. n.d. Google Scholar overview.

[10] Connected Papers. Connected Papers: Explore connected papers in a visual graph. n.d. connectedpapers.com.

[11] ResearchRabbit. ResearchRabbit: Literature discovery. n.d. researchrabbit.ai.

[12] University of Illinois. Illinois experts: Research and scholarship at UIUC. n.d. experts.illinois.edu.